

Grid Appliance – On the Design of Self-Organizing, Decentralized Grids

David Isaac Wolinsky, Arjun Prakash, Renato Figueiredo
University of Florida

Abstract—“Give a man a fish, feed him for a day. Teach a man to fish, feed him for a lifetime” – Lau Tzu

Grid computing projects such as TeraGrid [1], Grid’5000 [2], and OpenScience Grid [3] enable researchers to access vast amounts of compute resources, but in doing so, require researchers to adapt their workloads to the environments provided by these systems. Researchers do not have many alternatives as creating these types of systems involve coordination of distributed systems and expertise in networking, operating systems, security, and grid middleware. This results in many research groups creating small, in-house compute clusters where scheduling is often ad-hoc with low resource utilization. To address these challenges we present the “Grid Appliance.” The “Grid Appliance” enables researchers to efficiently deploy and extend their systems both locally and across network domains for both small and large scale computing grids. The paper details the design of the Grid Appliance and reports on experiences and lessons learned over four years of development and deployment of wide-area Grid appliance pools.

I. INTRODUCTION

Grid computing presents opportunities to combine various scale, distributed resources together to form powerful computing systems. Due to the challenges in coordinating the organization of grids, researchers typically become members of existing grids or otherwise often inefficiently managing their local resources, typically resource discovery and allocation by word of mouth. While there is a wealth of grid middleware available, such as resource managers like Condor [4], Torque (PBS) [5], and Sun Grid Engine [6] and parallelization tools like MPICH [7], Hadoop [8], and UPC [9], most researchers see the entry barrier to installing and management of the system as being greater than the usefulness of the application. Succinctly defined, we address the problem of how to introduce complex grid middleware to users, such that they can focus on making use of the software and not the setup and management of the systems thereafter.

Our approach begins with a P2P infrastructure enabling peers to discover each other coordinating the organization of a grid. This is performed through a distributed hash table (DHT) built on top of our P2P system. A DHT enables peers to efficiently query the system with a key and receive potentially multiple values. To address wide area connectivity and network dynamics, we use a virtual private network, guaranteeing all-to-all connectivity while transparently handling configuration and organization through the DHT. To avoid manual configuration of grid middleware, we have devised means to self-configuration through the DHT as well as IP multicast. The entire system has been packaged into a software

repository enabling automatic configuration of physical, virtual, and even cloud resources. The resulting “Grid Appliance” is a preconfigured environment emphasizing trivial installation. User interaction with the system is focused on the tools provided and not configuration details, providing researchers with a plug-and-play tool to create ad-hoc virtual compute clusters for their own groups, local or federated.

To allow user-friendly access to these feature, we have implemented a web interface called “Group Appliances.” At this site, users may create or join a group, similar to an online social networking group. As a member of the group, users can download a configuration file, which is used to configure the “Grid Appliance” and uniquely identify the user and group. Administrators of the group can allow users to join or remove misbehaving users. Once a user has obtained a configuration file, it can be added to the “Grid Appliance” as a floppy disk or into the file system. Upon the first boot, it will contact the “Group Appliances” site to obtain a certificate authority (CA) signed certificate and then connect to other systems through the P2P overlay. Afterward, the system is completely decentralized, relying on the P2P infrastructure for coordination.

To justify our techniques, consider the difficulty in combining resources across disparate networks, which may or may not involve multiple research groups. Challenges such as safety, security, connectivity, and efficiency may require an information technology (IT) expert. Network constraints present another complexity beyond configuration and organization of distributed resources. Contributing groups may have resources behind different network address translators (NATs) and firewalls, preventing direct communication with each other. Even assuming that an institution’s network administrator is willing to make exceptions for the grid, additional rules may be required for each new cluster or resource added internally and externally, quickly becoming unmanageable. Whereas in our approach, users can have a distributed system running by creating a group in a “Group Appliances” site, downloading the configuration image and a “Grid Appliance,” and starting the “Grid Appliance” virtual machine (VM).

The rest of this paper discusses these challenges in more depth and the solutions, we have developed, to address them. In Section II, we discuss various types of middleware and techniques used to support self-configuration. Section III reflects on the difficulties presented by distributed systems and how only a small subset of VPNs can provide a satisfactory solution. The entire architecture of the system is presented in

Section IV. Finally, Section V compares and contrasts other solutions to these problems.

II. MIDDLEWARE

We consider two types of grid middleware in this paper: resource management / batch task systems and application specific tools. Resource management systems like Torque, Condor, and Sun Grid Engine (SGE) consist of three fundamental components: execute nodes, resource managers, and submission nodes. The execute nodes run the tasks submitted from the submission sites. Users access a submission site, craft task description files, and submit them to a scheduler or resource manager, which will queue tasks to the various execute nodes. Application specific systems like Hadoop, MPICH, and UPC do not have such a convenient layout, but typically all-to-all connectivity is required amongst the resources.

A user configuring grid middleware must understand the layout of the system, install the correct software on each resource, and then configure all the individual resources to work with each other. Use of packaging as described in Section IV addresses the concerns of software. While the configuration data, users download from the “Group Appliances,” determines how to configure the individual resources. This section addresses decentralized resource configuration. Without it, users would, potentially, need to communicate with each other each time new resources were added, removed, or there was some system change, like an IP address changing. We explore the usage of both DHT and IP multicasting to allow users to forego this configuration issue.

A. Resource Managers and the DHT Approach

The first goal of the “Grid Appliance” is to construct wide area grids using common resource management techniques. To do so efficiently, we employ the DHT. In this manner, a manager places its IP address into the DHT at the key *managers*. When workers and clients join the grid, the systems automatically query this key and configure to report to one or more managers, application dependent. Likewise, managers can query this key and then coordinate with each other management of the overlay.

Of the resource management middlewares that we have surveyed, Condor was the most appealing due to its decentralized properties and focus on desktop grids. These are easily recognizable by analyzing the complexity involved in having multiple submission sites and the ability to add and remove resources.

Typical clusters usually have a single entry point for task submission, SGE and Torque make it rather difficult to support more than one, usually requiring additional middleware, such as Globus [10], to support this behavior. A single submission site may be very undesirable and difficult to support, if multiple groups are collaborating together to form a single grid. To allow for multiple submission sites, Condor separates job queues / submission sites from resource management and so that two entities work together to negotiate the running jobs

on resources. In our system, job queue or submission machine, which stores everything pertinent to the user, is run by the user, while the resource manager runs independently with Condor handling all the interaction between it and the job queue.

To add new resources to a Condor system, an execute or submission node must have an IP address for the manager, the rest of the system organization is performed entirely transparent to the user. However, in SGE and Torque, after resources have been added into the system, the user must manually configure the manager to control them.

One of the caveats of our approach, which remains ongoing research, is the requirement of a manager, which means we do not provide a completely decentralized, but rather a distribute approach. In the meantime, we have taken advantage of a feature known as “flocking” in Condor. Flocking allows submission sites to connect to multiple managers. This serves two purposes: 1) to provide transparent reliability by supporting multiple managers and 2) users can share their resources through their own manager.

To configure Condor, we store managers IP addresses into the DHT using the key *managers*. When a new peer joins, it queries the DHT, obtains the list of all managers, and randomly selects one as its primary manager. The rest are set to flocking. If the system prefers managers from its group, it will randomly contact each manager in an attempt to find a match, selecting one at random if no match is found. If no managers are found, the process repeats every 60 seconds. Once a manager has been found, it is checked every 10 minutes to ensure it is online and additional managers that have come online are added to the flock list.

B. Hadoop, MPICH, and Multicast Discovery

Alternatively users may want to experiment with tools meant for LANs but not want to invest the time to install and configure them. In which case, DHT use does not translate well if they want to install software without using our appliance domain. To support these endeavours, we have investigated methods to bootstrap grid middleware through IP multicast. IP multicast works on all LANs and is used by many popular applications for discovery, such as Windows Media Center and iTunes by means of the UPNP (DLNA) and DAAP protocols, respectively.

The two systems that we have used to configure through IP multicast are Hadoop and MPICH. Hadoop configuration consists of a head node and worker nodes, where the head node distributes map tasks to the worker nodes. In MPICH, each resource is identically configured to support interprocess communication through the MPICH library, it is up to the application developer to determine roles of individual nodes. In both cases, the way multicast discovery works is to send out a beacon requesting that all nodes supporting these services respond.

Our appliance setup for these applications consists of a common ssh key, so that all resources in a grid can connect with each other and a multicast discovery application so a coordinator can organize the resources. The ssh key only




Fig. 1. A typical NAT interaction. The peer behind a NAT has a private address. When the packet is sent through the NAT, the NAT translates the source information into a public mapping, keeping the original source information so that if a packet from the remote peer comes back, it can be translated and delivered to the original source.

enables access on host only networks and VPNs, preventing malicious third-parties from gaining access to the grids. The IP multicast script is run by the user on a single access node, which will act as the coordinator. The resource discovery sends out a beacon several times and after a 30 second delay, the process completes and all responding resources are automatically configured through ssh.

III. THE MOTIVATION FOR VPNs

As of 2010, the majority of the Internet is connected via Internet Protocol (IP) version 4. This protocol has a quickly approached limit of addresses available, only 2^{32} (approximately 4 billion). With the Earth's population at over 6.8 billion and each individual potentially having multiple devices with Internet connectivity, the IPv4 limitation is becoming more and more apparent. Addressing this issue are two approaches: 1) the use of NATs to enable many machines and devices to share a single IP address but preventing bidirectional connection initiation, and 2) IPv6 which supports 2^{128} addresses. The use of NATs, as shown in Figure 1, complicates grid systems that require all-to-all communication, which include all of those which we consider. In addition, firewalls may prevent peers from receiving incoming connections. And while the eventual widespread use IPv6 may eliminate the need for address translation, it does not deal with the issue of firewalls, and the future of NATs in IPv6 is unclear.

VPNs motivate from more than just crossing NATs and firewalls. With a VPN, users can avoid the headaches associated with dynamic IP addresses, as each resource can claim an IP and, regardless of the machines physical location, use it, an ideal condition for laptop users. In addition, it abstracts the user from having to be concerned about network addresses. The mixture of virtual machine NATs and VPNs allow users inside a network to not be concerned about IP address allocations, from the LANs perspective, the machine has only ever been allocated one IP address, yet all grid resources are able to communicate with each other.

Our work relies on a group enabled IPOP VPN called GroupVPN [11]. IPOP through its underlying P2P infrastructure support NAT traversal allowing for peers behind NATs to communicate directly with each other in addition to indirectly communicating by relaying across the P2P system. Many of the key features of our grid system are enabled because of the VPN. For example, if there were not a VPN, users of

MPI and Hadoop would need to ensure that all resources were bridged to the LAN and not through a VM NAT, the typical configuration, otherwise the multicast message would not be delivered to all participants. The VPN software supports the ability to self-organize using existing infrastructures including IP multicast, public overlays, and Xmpp as described in our previous work [12]. This is in contrast to other VPNs that are either centralized and require a dedicated node to coordinate peers and decentralized solutions that require manual configuration of links between peers.

Using these techniques "Grid Appliances" can be constructed in one of two ways: local and wide area. The "Grid Appliance" ships with two default configurations, one that connects users to a globally available public system and another that allows for LAN only grids. Local grids can be constructed by booting the appliances, which will then use multicast self-discovery to find other resources, create the DHT overlay, and then form VPN links. Alternatively, the user can connect to the default public system or use "Group Appliances" to create and manage their own grid, both of which bootstrap from a public shared DHT overlay. This does not prohibit more advanced users from downloading our "Group Appliances," as its available as a VM, and host their own DHT overlay.

IV. DEPLOYING GRID APPLIANCES

The complete system is presented in Figure 2. Users join a group and can then download one of three different types of configurations: server, worker, and client (mapping to the three types of typical grid nodes, resource manager, execution site, and submission site). In addition to the configuration data, a first time user can download the generic appliance, or one extended with additional software customized for their community. In the case of a VM appliance, the configuration data can be added to the system as a virtual floppy image, automatically configuring the appliance for a specific behavior. In the case of the client, it begins searching for an appropriate resource manager by querying the DHT. Upon finding an entry in the DHT, the client communicates with the resource manager via the VPN. The same procedure can be repeated many times to add additional resources into the grid. Resources booted using the "server" image configure and advertise themselves through the DHT or VPN. Each individual grid can support multiple servers for fail-over or

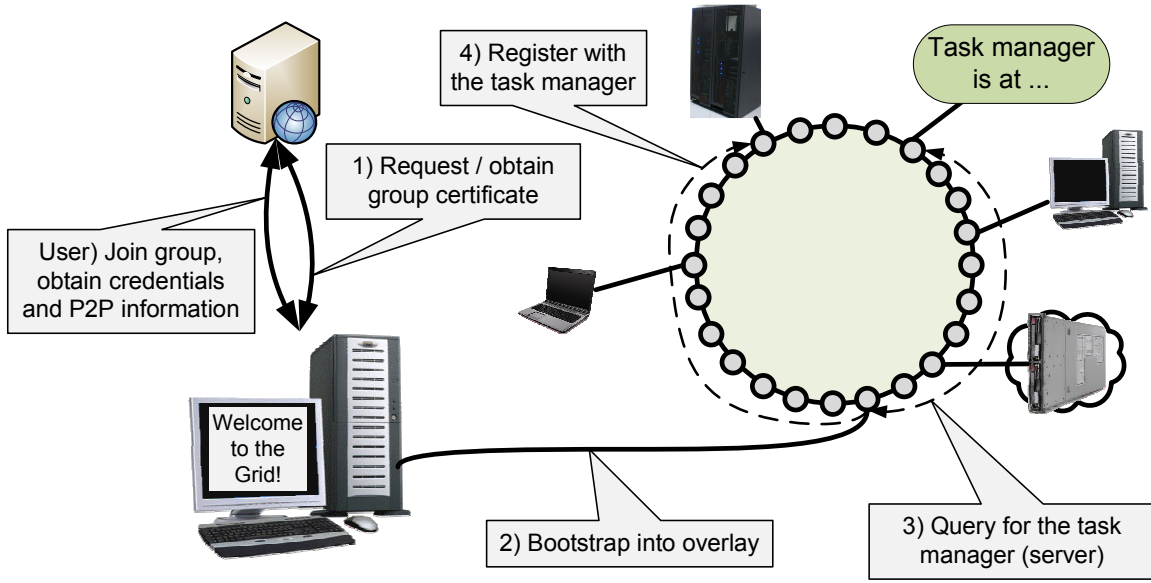


Fig. 2. An example deployment scenario: obtaining configuration files, to starting the appliance, and connecting with a resource manager.

load balancing, depending on the capabilities of the grid middleware (e.g. Condor flocking). Beyond the downloading of the appliance and the configuration data, all other steps are performed transparently from the user.

Often VMs are favored for the distribution of complicated applications as experts can configure them and release the results as a complete working system. This approach may limit non-expert use to the VM appliance, which may be undesirable for users that want to configure their own systems without reuse of the existing VM. Guides may exist for the creation of systems, most systems are too complex for non-experts to produce similar results found in the VM. In addition to supplying VMs, we also supply packages for Debian and Ubuntu systems, DEB files. These enable easy installation in arbitrary environments and through the use of package managers, APT in the case of Debian and Ubuntu, which handle configuration such users can focus on the end result and not stumble during setup. Our packages install a complete working system given a working installation of Ubuntu or Debian, such that the a fresh installation of Ubuntu whether on a physical machine or on the cloud can be configured by following these steps: add the configuration data into the newly configure resource, our package repository to the managers list, and our repositories signing key to the key chain for repositories and finally selecting our packages and installing them. At which point, the system will connect to the grid.

To verify the utility of our approach, we have evaluated the time required to create and utilize a grid consisting of various distributed resources using the “Grid Appliance.” The system consists of VMware resources behind a Cisco and “iptables” NAT at the University of Florida (UF), KVM resources behind an “iptables” and KVM NAT at Northeastern University (NEU), and cloud resources provided by Amazon’s EC2, pools

of 50 resources from each site were booted independently and then together, resulting in 4 different test runs.

Each test was performed independently with an existing manager and submit node already in the system. We measure three times during this evaluation: “start” - begins with starting the experiment including the copying of files and creation of instances until all resources have been powered on, “connect” - is the time delta between the end of “start” and when all resources appear in “condor_status,” that is, have registered with the manager, and “run” - time from the submission to the conclusion of a 5 minute job to all resources. Like connect, run measures the time for VPN connections, only from the client to the resources instead of from the negotiator. All tasks are automated through scripts with human interaction required only to start the events of grid boot and job submission. Results are presented in Figure 3.

	50 - EC2	50 - NEU	50 - UF	150 - All
Start	2:44	10:21	20:23	21:14
Connect	2:27	11:36	3:53	17:13
Run	7:15	6:35	5:53	21.19

Fig. 3. Time in minutes:seconds to start and connect execute nodes from various sites, Amazon EC2, Northeastern University, and University of Florida, to an already online resource manager, and then the time to run a 5 minute job from a freshly connected submission node.

As the systems consist of various hardware and software configurations, the time to start is only provided as a reference to potential overheads in bootstrapping the resources. Some of the interesting experiences of the experiment were: 1) the combination of the “iptables” and VMware NAT was more easily traversable than the combination of “iptables” and KVM NAT and 2) in the experiment consisting of 150 peers, nodes were actually well connected much earlier, but due to missed

packets and Condor timeouts, not all resources were accounted for in Condor as early as in the other tests. With regards to the KVM NAT, it appears to be particularly aggressive as NAT mappings last for less than 10 seconds, while typical NATs keep mappings for over 30 seconds.

V. RELATED WORK

Existing work that falls under the general area of desktop grids/opportunistic computing include Boinc [13], BonjourGrid [14], and PVC [15]. Boinc, used by many “@home” solutions, focuses on adding execute nodes easy; however, job submission and management rely on centralization and all tasks must use the Boinc APIs. BonjourGrid removes the need for centralization through the use of multicast resource discovery; the need for which limits its applicability to local area networks. PVC enables distributed, wide-area systems with decentralized job submission and execution through the use of VPNs, but relies on centralized VPN and resource management.

Each approach addresses a unique challenge in grid computing, but none addresses the challenge presented as a whole: easily constructing distributed, cross-domain grids. Challenges that we consider in the design of our system are ensuring that submission sites can exist any where and are not confined to complex configuration or highly available, centralized locations; ability to dynamically add and remove resources by starting and stopping an appliance; and the ability for individual sites to share a common server or to have one or more per site so that no group in the grid is dependent on another. We emphasize these points, while still retaining the ease of use of Boinc, the connectivity of PVC, and the flexibility of BonjourGrid. The end result is a system similar in organization to OurGrid [16], though with self-configuration and organization and a VPN to transparently handle network constraints.

VI. CONCLUSIONS

In this paper, we have described a novel grid architecture that enables both wide area and educational grid middleware. Our results show that the process of connecting the resources together are not significantly longer than that of actually starting the resources and that submission sites have very low overheads in connecting to the resources. Furthermore, we have practical experience of using the system for over 2 years through a project called Archer [17], an active grid deployed for computer architecture research. Archer currently spans four seed universities with 500 resources with over hundreds of students and researchers submitting jobs totaling over 150,000 hours of job execution in the past year alone. Groups at the Universities of Florida, Clemson, Arkansas, and Northwestern Switzerland have used it as a tool to teach grid computing. Clemson and Purdue are constructing campus grids using the underlying VPN to connect resources together. Recently, two private small-scale systems have come online using our shared system available at www.grid-appliance.org. Feedback from users through surveys have shown that non-expert users are

able to connect to our public Grid appliance pool in a matter of minutes by simply downloading and booting a plug-and-play VM image that is portable across VMware, VirtualBox, and KVM.

REFERENCES

- [1] C. Catlett, “The philosophy of teragrid: Building an open, extensible, distributed terascale facility,” in *CCGRID '02: Proceedings of the 2nd IEEE/ACM International Symposium on Cluster Computing and the Grid*. Washington, DC, USA: IEEE Computer Society, 2002, p. 8.
- [2] F. Cappello, E. Caron, M. Dayde, F. Desprez, E. Jeannot, Y. Jegou, S. Lanteri, J. Leduc, N. Melab, G. Mornet, R. Namyst, P. Primet, and O. Richard, “Grid’5000: a large scale, reconfigurable, controllable and monitorable Grid platform,” in *Grid’2005 Workshop*. Seattle, USA: IEEE/ACM, November 13-14 2005.
- [3] R. Pordes, D. Petravick, B. Kramer, D. Olson, M. Livny, A. Roy, P. Avery, K. Blackburn, T. Wenaus, F. Wrthwein, I. Foster, R. Gardner, M. Wilde, A. Blatecky, J. McGee, and R. Quick, “The open science grid,” vol. 78, no. 1, 2007, p. 012057.
- [4] M. Livny, J. Basney, R. Raman, and T. Tannenbaum, “Mechanisms for high throughput computing,” *SPEEDUP Journal*, June 1997.
- [5] Cluster Resources. (2010, July) Torque resource manager. <http://www.clusterresources.com/pages/products/torque-resource-manager.%php>.
- [6] Sun. (2010, July) gridengine. <http://gridengine.sunsource.net/>.
- [7] A. N. Laboratory. (2010, July) MPICH2. <http://www.mcs.anl.gov/research/projects/mpich2/>.
- [8] Apache. Hadoop.
- [9] T. El-Ghazawi, W. W. Carlson, and J. M. Draper. (2003, October) UPC language specification v1.1.1.
- [10] Globus Alliance. (2010, July) Globus toolkit. <http://www.globus.org/toolkit/>.
- [11] D. I. Wolinsky and et al., “On the design and implementation of structured P2P VPNs,” in *ARXIV 1001.2575*, 2010.
- [12] D. I. Wolinsky, P. St. Juste, P. O. Boykin, and R. Figueiredo, “Addressing the P2P bootstrap problem for small overlay networks,” in *10th IEEE International Conference on Peer-to-Peer Computing*, 2010.
- [13] D. P. Anderson, “Boinc: A system for public-resource computing and storage,” in *the International Workshop on Grid Computing*, 2004.
- [14] H. Abbes, C. Cérin, and M. Jemni, “Bonjourgrid: Orchestration of multi-instances of grid middlewares on institutional desktop grids,” in *IPDPS*, 2009, pp. 1–8.
- [15] A. Rezmerita, T. Morlier, V. Neri, and F. Cappello, “Private virtual cluster: Infrastructure and protocol for instant grids,” in *Euro-Par*, 2006.
- [16] N. Andrade, L. Costa, G. Germoglio, and W. Cirne, “Peer-to-peer grid computing with the ourgrid community,” in *23rd Brazilian Symposium on Computer Networks (SBRC 2005) - 4th Special Tools Session*, 2005.
- [17] R. J. Figueiredo, P. O. Boykin, J. A. B. Fortes, T. Li, J. Peir, D. Wolinsky, L. K. John, D. R. Kaeli, D. J. Lilja, S. A. McKee, G. Memik, A. Roy, and G. S. Tyson, “Archer: A community distributed computing infrastructure for computer architecture research and education,” in *Collaborative Computing: Networking, Applications and Worksharing*, vol. 10. Springer Berlin Heidelberg, 2009, pp. 70–84. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-03354-4_7